

Optimizing Unity Games

Aleksandr Dolbilov

CTO, JoyCraft Games



About me

Aleksandr Dolbilov

- 10+ years in game industry
- Android, iOS, PS2, PS3, PSP, XBox, XBox 360...



Nitro Nation



Agenda

- Performance problems from real-world games
 - Live profile, bottleneck detection and optimization
 - Unity profiler and platform-specific tools in action
 - Solutions and tips which help to improve performance of your games
-
- **Grab your projects today and perform live performance analysis tomorrow!!**



Today

- Few notes about performance in games
- CPU optimization
 - Static scene
 - Vertex constants arrays
 - Vertex constants instancing
 - Math optimization



Tomorrow...

- GPU performance optimization
 - Vertex shaders optimization
 - Fragment shaders optimization





Performance in games



Frame time



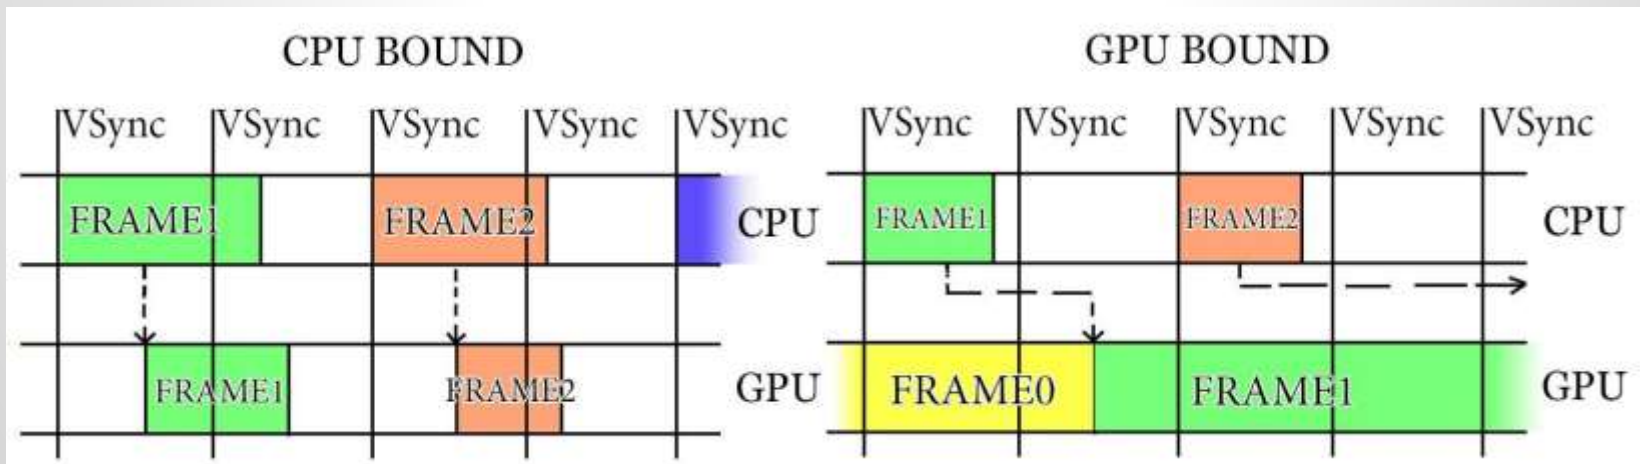
- Frame time depends on **CPU** time and **GPU** time
- Both units work in parallel if you don't touch **GPU** data from **CPU**



Locate your bottleneck



- Detect if game is **CPU** or **GPU** bound
 - ***Gfx.WaitForPresent*** in Unity Profiler is good indicator
 - Platform-specific tools provide units utilization info



Locate your bottleneck (CPU bound)

- Use **Unity Profiler** for further analysis
- Use ***Profiler.BeginSample*** and ***Profiler.EndSample*** to get detailed information on device
- **Deep profile** option is very helpful to find hidden pitfalls (implicit type casts, redundant constructors etc.)
- **Time.realtimeSinceStartup** can help you to monitor performance for release builds



Locate your bottleneck (GPU bound)

- Use platform-specific tools

- *Adreno Profiler* (Qualcomm, Adreno)
- *PerfHUD ES* (nVidia, Tegra)
- *PVRTune, PVRTrace* (Imagination tec., PowerVR)
- *Mali Graphics Debugger* (ARM, Mali)

- Show/hide method



Optimization tips

- Always measure optimization results
 - Use milliseconds (not FPS!!)
 - If it's possible, create simple test for faster iterations
-
- Check if everything work properly 😊



CPU performance optimization



Render CPU time measurement

- Check ***Camera.Render*** time in Unity Profiler
- Or create custom ***MonoBehaviour***
 - ***OnPreCull*** is timestamp begin
 - ***OnPostRender*** is timestamp end
 - ***OnGUI*** for render result



Static scene



Initial scene

GameObject count: 1440
Material count: 32
Shader count: 2
Textures count: 32
Draw calls: 1434




CPU time (HTC One S): ~75 ms



Combine geometry: how to



- Usually made by artists according to 'common sense'



Combine geometry: motivation

- Reduce engine overhead on **Transform** and **MeshRenderer** processing
 - Geometry transform validation
 - Visibility/Occlusion culling
 -
- Reduce draw calls count



Combine geometry: result

GameObject count: **51**/1440
Material count: 32
Shader count: 2
Textures count: 32
Draw calls: **90**/1434



CPU time (HTC One S): ~**5.2 ms** / ~75 ms



Bake textures to atlases: how to

- Can be made manually by artists but...
- ... better to automate the process with script (use ***Texture2D.PackTextures***)



Bake textures to atlases: how to

- Our approach is to create separate atlases for opaque and transparent objects
- Tiled textures are excluded from atlas
- Huge textures are excluded from atlas
- 2048x2048 texture resolution is usually enough to fit all static scene textures



Bake textures to atlases: motivation

- Reduce **Material** and **Texture** count
- Reduce engine overhead on **Material** setup
- Reduce engine overhead on **Texture** switch



Bake textures to atlases: result

GameObject count: 51
Material count: 4/32
Shader count: 2
Textures count: 4/32
Draw calls: 90



CPU time (HTC One S): ~**3.2 ms**/~5.2 ms



Move geometry to world space

- Pre-transform geometry to world space
- Use **Shader.SetGlobalVector** for ViewProjection matrix:
 - Reduce engine overhead on $\text{World} * \text{ViewProjection}$
 - Reduce engine overhead on shader constants setup, because Unity has setup cache for Vector properties
- We have ~25% performance boost on some scenes
- CPU time (HTC One S): ~**2.8 ms**/~3.2 ms



Static batching

- Reduce draw calls count but...
- ... this is not for free
 - Performance depends on triangle count
 - Some devices doesn't like dynamic geometry
 - We have severe performance degrade on **Mali** devices
 - Doesn't get any speed improvement for **HTC One S**
 - Your mileage may vary




Vertex shader arrays



Motivation

- Reduce draw call count for objects with slightly different materials



General algorithm

- Combine all meshes and store **MaterialID** for each vertex (index in array)
- Add array of uniforms to vertex shader
- Add value extraction code to vertex shader
- Map array elements as material properties



MaterialID pack/unpack

- Using texcoord (uv/uv2)

- pack: ***mesh.uv[i].x = mesh.uv[i].x + MaterialID;***
- unpack: ***int MaterialID = (int) v.texcoord.x;***

- Using color

- pack: ***mesh.colors32[i].a = MaterialID;***
- unpack: ***int MaterialID = (int) (v.color.a*255.0f + 0.5f);***



Shader code

```
CGPROGRAM
```

```
...
```

```
float4      _Colors[9];
```

```
...
```

```
int         materialID = (int) v.texcoord.x;
```

```
float4      _Color      = _Colors[materialID];
```

```
...
```

```
ENDCG
```



Material properties

- Each array element can be mapped as individual property

PropertyName = ArrayName+ElementIndex

```
_Colors0 ("Color doors", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors1 ("Color front bumper", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors2 ("Color rear bumper", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors3 ("Color hood", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors4 ("Color base", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors5 ("Color mirrors", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors6 ("Color spoiler", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors7 ("Color top", Color) = (1.0, 1.0, 1.0, 1.0)
_Colors8 ("Color trunk", Color) = (1.0, 1.0, 1.0, 1.0)
```

Usage tips

- Don't make vertex array too large
 - Every array element require **glUniform...** call (in Unity3D)
 - It can significantly degrade vertex shader performance on some platforms
- You can use custom property names to display in inspector
 - hood color, body color, etc..



Results

	Before	After
Draw calls	37	29
Frame time (HTC One S), ms	3.7	2.7



Vertex constants instancing



Motivation

- Reduce draw call count for objects with similar materials even if they have different transform




General algorithm

- Create custom Mesh which contains N copies of original Mesh with InstanceID stored in vertex
 - Set custom bounding box
 - Add array of uniform matrices to vertex shader
- (InstanceData)**
- Add value extraction code to vertex shader
 - Write custom script to update **InstanceData**



Visibility problem

- Create custom bounding box which preserve geometry to become hidden when it should be visible





Custom setup code

- Use **OnWillRenderObject** to setup your data
 - called only if object is visible
- Use **propertyID** for material setup
 - works much faster than usual 'string' version

```
void OnWillRenderObject () {  
    material.SetMatrix(propertyIDs[0], _transforms[0].localToWorldMatrix);  
    material.SetMatrix(propertyIDs[1], _transforms[1].localToWorldMatrix);  
    material.SetMatrix(propertyIDs[2], _transforms[2].localToWorldMatrix);  
    material.SetMatrix(propertyIDs[3], _transforms[3].localToWorldMatrix);  
}
```

Results

	Before	After
Draw calls	29	17
Frame time (HTC One S), ms	2.7	2.0



Efficiency depending on batch size

Efficiency = timeBefore/timeAfter

	4 elements batch	8 elements batch	16 elements batch
Mali-400 MP	x2.8	x4.0	x6.0
Adreno 220	x2.9	x5.0	x6.5
PowerVR SGX 540	x3.1	x5.0	x6.5
nVidia Tegra 3	x2.5	x4.0	x5.0

Math optimization



Case study

- Getting 20 000 matrices which transform object from local space to camera space
- Naive implementation: **125 ms** (HTC One S)

```
public static void ApplyTransform(Matrix4x4[] outputMatrices, Matrix4x4[] inputMatrices)
{
    for(int i = 0; i < inputMatrices.Length; ++i) {
        outputMatrices[i] = Camera.main.worldToCameraMatrix*inputMatrices[i];
    }
}
```

Cache complex expressions



- Most of the properties evaluate complex expression every time they invoked!! (**Transform.forward**, **Vector3.zero** ...)
- Sometimes they result to safe/unsafe context switch
- Optimized implementation: **33.5 ms**/125 ms (HTC One S)

```
Matrix4x4    worldToCameraMatrix = Camera.main.worldToCameraMatrix;
for(int i = 0; i < inputMatrices.Length; ++i) {
    outputMatrices[i] = worldToCameraMatrix*inputMatrices[i];
}
```



Remove redundant copies

```
public static Matrix4x4 operator *(Matrix4x4 lhs, Matrix4x4 rhs)
```

- Matrix4x4 is value-type
 - We have 2 redundant copies for input arguments
 - We have 1 redundant copy for output result
- Create our own method using reference semantic
- Optimized implementation: **21.5 ms**/33.5 ms (HTC One S)

```
static void MultiplyMatrices(ref Matrix4x4 result, ref Matrix4x4 lhs, ref Matrix4x4 rhs)
```

Remove redundant calculations

- All our matrices has **Matrix43** semantic

$$M = \begin{pmatrix} m00, & m01, & m02, & m03 \\ m10, & m11, & m12, & m13 \\ m20, & m21, & m22, & m23 \\ (0, & 0, & 0, & 1) \end{pmatrix}$$

- We can remove a lot of calculations using this fact
- Optimized implementation: **15.2 ms**/21.5 ms (HTC One S)



Write native plug-in (C code)

- You can use functions from android shared libraries using Platform Invoke Service
- Create your own library using Android-NDK
- Create functions with C-linkage (***extern "C"*** for *.cpp)
- Define your own ***struct Matrix4x4*** (be careful with elements order !!)
- Pass and return elements using pointers
- Now you can use all C++ optimizations !!



Write native plug-in (C code)



- Simply port our optimized operator to C

```
struct Matrix4x4 {  
    float m22; float m32; float m02; float m12;  
    float m23; float m33; float m03; float m13;  
    float m20; float m30; float m00; float m10;  
    float m21; float m31; float m01; float m11;  
};  
  
extern "C" void MultiplyMatricesArrayC(Matrix4x4* outputMatrices,  
                                       const Matrix4x4* worldToCamera,  
                                       const Matrix4x4* inputMatrices,  
                                       int count)
```

Write native plug-in (C# code)



- Insert '***using System.Runtime.InteropServices;***'
- Declare P/Invoke method

```
#if !UNITY_EDITOR && UNITY_ANDROID
    [DllImport("NativePlugin")]
    private static extern void MultiplyMatricesArrayC(
        ref Matrix4x4      outputMatrices,
        [In] ref Matrix4x4 worldToCamera,
        [In] ref Matrix4x4 inputMatrices,
        int                count
    );
#endif
```


Write native plug-in (C# code)

- Pass pointers to array as reference to first element

```
#if !UNITY_EDITOR && UNITY_ANDROID  
    MultiplyMatricesArrayC(ref outputMatrices[0],  
                           ref worldToCameraMatrix,  
                           ref inputMatrices[0],  
                           inputMatrices.Length);  
#else
```



Write native plug-in

- Native plug-ins has constant overhead due to safe/unsafe context switch, so try to use them for large batches of data
 - Make sure you enable all compiler/linker optimizations to get maximum performance
-
- Optimized implementation: **6.8 ms**/15.2 ms (HTC One S)



Use ASM code



```
#include <arm_neon.h>

void __attribute__((noinline))
_MultiplyMatricesArrayNeon(
    float* outputMatrices,
    const float* worldToCamera,
    const float* inputMatrices,
    int count
)
{
    count -= 2;

    asm volatile
    (
        "pld [%input]                \n\t"
        "vld1.32 {d8-d11}, [%w2c]!   \n\t"
        "vld1.32 {d12-d15}, [%w2c]   \n\t"

        // prologue
        "vld1.32 {d0-d3}, [%input]!   \n\t"
        "vmul.f32    q9, q4, d2[0]     \n\t"
        "vmul.f32    q8, q4, d0[0]     \n\t"
        "vmla.f32    q9, q5, d2[1]     \n\t"
        "vmla.f32    q8, q6, d1[0]     \n\t"
        "vmla.f32    q9, q6, d3[0]     \n\t"
        "vmla.f32    q8, q7, d1[1]     \n\t"
        "vmla.f32    q9, q7, d3[1]     \n\t"
        "vld1.32 {d4-d7}, [%input]!   \n\t"

        // main loop
        "Lloop2:                      \n\t"
        "vld1.32 {d0-d3}, [%input]!   \n\t"
        "                                \n\t"
        "subs        %[count], %[count], $1 \n\t"
        "bgt         Lloop2           \n\t"

        // prefetch
        "pld [%input], #1024          \n\t"

        "vmul.f32    q10, q4, d4[0]    \n\t"
        "vmul.f32    q11, q4, d6[0]    \n\t"
        "vmla.f32    q10, q6, d5[0]    \n\t"
        "vmla.f32    q11, q6, d7[0]    \n\t"
        "vmla.f32    q10, q7, d5[1]    \n\t"
        "vmla.f32    q11, q7, d7[1]    \n\t"
        "vst1.32 {d16-d19}, [%output]! \n\t"

        "vld1.32 {d4-d7}, [%input]!   \n\t"
        "vmul.f32    q9, q4, d2[0]     \n\t"
        "vmul.f32    q8, q4, d0[0]     \n\t"
        "vmla.f32    q9, q5, d2[1]     \n\t"
        "vmla.f32    q8, q6, d1[0]     \n\t"
        "vmla.f32    q9, q6, d3[0]     \n\t"
        "vmla.f32    q8, q7, d1[1]     \n\t"
        "vmla.f32    q9, q7, d3[1]     \n\t"
        "vst1.32 {d20-d23}, [%output]! \n\t"

        // epilogue
        "vmul.f32    q10, q4, d4[0]    \n\t"
        "vmul.f32    q11, q4, d6[0]    \n\t"
        "vmla.f32    q10, q6, d5[0]    \n\t"
        "vmla.f32    q11, q6, d7[0]    \n\t"
        "vmla.f32    q10, q7, d5[1]    \n\t"
        "vmla.f32    q11, q7, d7[1]    \n\t"
        "vst1.32 {d16-d19}, [%output]! \n\t"
        "vst1.32 {d20-d23}, [%output]! \n\t"

        :
        : [output]          "r"(outputMatrices)
        : [input]           "r"(inputMatrices)
        : [w2c]             "r"(worldToCamera)
        : [count]           "r"(count)
        : "cc", "memory"
        : "q0", "q1", "q2", "q3", "q4", "q5"
        : "q6", "q7", "q8", "q9", "q10", "q11"
    );
}
```

Use ASM code

- Lowest possible level
 - cache prefetch
 - parallel load/store
 - hardware registers management
 - ...
- NEON extension is really helpful for vector/matrix math
- Optimized implementation: **3.3 ms**/6.8 ms (HTC One S)



Don't give up 😊

- Final result **~40 times faster** than naive approach !!
- For most situations it's enough to perform smart C# optimizations (**~8 times faster**)
- But keep in mind native plug-ins for computationally intensive tasks (additional **~5 times faster**)



GPU performance optimization



Mobile GPU vendors

- Imagination Technologies (PowerVR)
 - Series 5, 5XT, 6, 6XT, Wizard
- ARM (Mali)
 - Utgard (Mali 400), Midgard (Mali T624)
- Qualcomm (Adreno)
 - Adreno 2xx, 3xx, 4xx
- nVidia (Tegra)
 - Tegra 2, 3, 4, K1



Rendering approach PowerVR

- Tile-Based Deferred Rendering (TBDR)
 - pixel-perfect HSR (Hidden Surface Removal)
 - small tiles 32x32 (on-chip registers)
- Scalar ALU architecture (from 6 series)
- Different fragment shader precision supported
- Unified ALU
- High-quality fast MSAA
- 32-bit internal color precision



Rendering approach Mali

- Tile-Based Rendering (TBR)
 - early-Z used for HSR
 - small tiles 16x16 (on-chip registers)
- Separate ALU (Utgard), unified ALU (Midgard)
- Different fragment shader precision supported
- High-quality fast MSAA
- 32-bit internal color precision



Rendering approach Adreno

- Tile-Based Rendering (TBR)
 - early-Z used for HSR
 - large tiles (about 256 Kb, on-chip GMem)
 - binning algorithm to classify polygons
 - vertex shader can be called multiple times
- Unified ALU
- Different fragment shader precision supported
- Scalar ALU (from 3xx)
- MSAA can significantly increase vertex processing cost
- Color precision can be used to reduce tile count



Rendering approach Tegra

- Immediate-Mode Rendering
 - early-Z used for HSR (Hi-Z for Tegra K1)
 - render directly to framebuffer
 - framebuffer compression (from Tegra 4)
- Separate ALU (Tegra 2, 3, 4), unified (Tegra K1)
- CSAA (Tegra 2, 3)
- MSAA (Tegra 4, K1)
- Color precision can reduce bandwidth



Vertex shaders optimization



Motivation

- High-poly meshes
- Some **GPUs** have separate vertex **ALU**
- Adreno **GPU** can process vertices multiple times due to binning algorithm (especially when **MSAA** is enabled)



Optimize before vertex ALU

- Always enable '**Optimize Mesh**' option in Mesh Import Settings
 - Post-transform vertex cache optimization
 - Pre-transform vertex cache optimization
- Always enable '**Optimize Mesh Data**' option in '**Player Settings->Other Settings**'
 - Remove redundant vertex attributes (tangents, normals, color etc.)



Case study: Car paint vertex shader

- Use a lot of features
 - diffuse light from 3 light sources
 - fresnel
 - tunnel mapping
 -



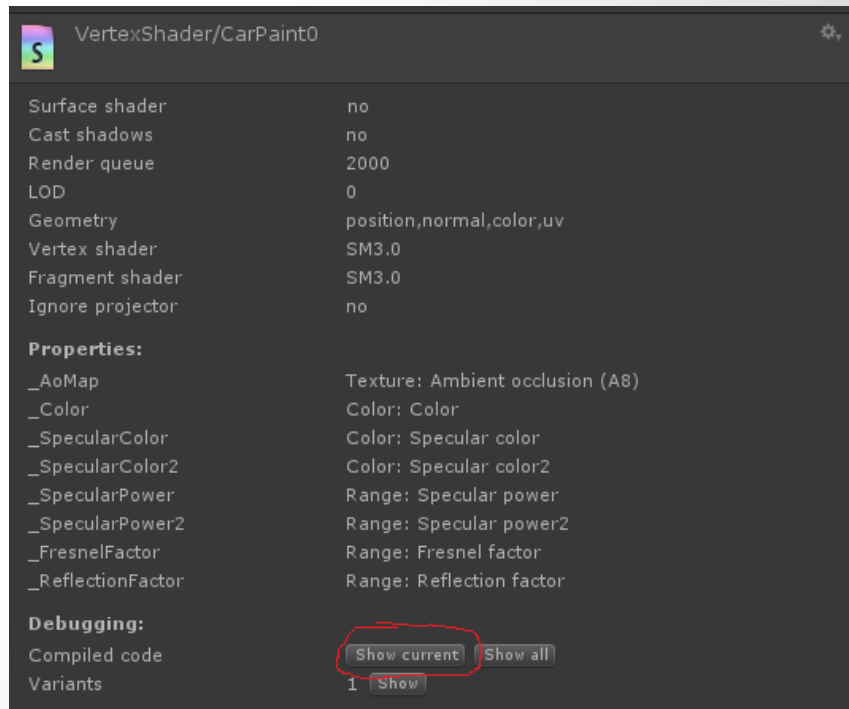
- Optimize for **PowerVR SGX 540**



Optimization pipeline



- Grab **GLSL ES** code



Optimization pipeline



- Copy code from '**#ifdef VERTEX**' section

```
SubProgram "gles " {
    "!!GLSL

#ifdef VERTEX

attribute vec4 _glesVertex;
attribute vec3 _glesNormal;
attribute vec4 _glesMultiTexCoord0;
uniform highp vec3 _WorldSpaceCameraPos;
uniform highp mat4 glstate_matrix_mvp;
uniform highp mat4 _Object2World;
uniform lowp vec4 _Color;
uniform highp float _FresnelFactor;
uniform highp float _ReflectionFactor;
uniform highp float global_TunnelAxisScale;
uniform highp float global_TunnelAxisOffset;
uniform highp float global_TunnelAxisReflectionScale;
uniform highp float global_TunnelAngleReflectionScale;
uniform highp vec4 global_LightDirection1;
uniform highp vec4 global_LightColor1;
uniform highp vec4 global_LightDirection2;
uniform highp vec4 global_LightColor2;
uniform highp vec4 global_LightDirection3;
uniform highp vec4 global_LightColor3;
uniform highp vec4 global_Ambient;
varying mediump vec2 xlv_TEXCOORD0;
varying mediump vec2 xlv_TEXCOORD1;
varying lowp vec4 xlv_TEXCOORD2;
varying mediump vec3 xlv_TEXCOORD3;
void main ()
{
    highp vec3 normalizedReflectedViewVectorInWorldSpace_1;
    lowp vec3 diffuse_2;
    highp vec4 tmpvar_3;
```

Optimization pipeline



- Paste code into **PVRShaderEditor**
- Select appropriate compiler '**Preferences->Compiler Settings->Compiler**'
- Analyze the result

Per-Line Cycle Estimate Total: 105		Emulated Cycle Total: 88-98	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	88		
Emulated Cycles (Worst Case):	98		
Temporary Registers Used:	7		
Primary Attributes Used:	11		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	3		
<hr/>			
SL:ES Vertex Shader	Line: 30	Col: 25	INS

Analyze the result

- Emulated Cycles (Worst Case)
 - maximum number of cycles spent by **ALU** to process single vertex
- Temporary Registers Used
 - **ALU** has limited amount of temporary registers
 - the number of shader threads that work simultaneously depends on this parameter
 - more threads help you to avoid stalls when vertex attribute fetch occurs



Bake complex calculations



- ***atan2*** function consume too many cycles
- We can pre-compute it using ***AssetPostprocessor.OnPostprocessModel***
- Store result in ***mesh.uv2.x***

Per-Line Cycle Estimate Total: 80		Emulated Cycle Total: 70	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	70		
Temporary Registers Used:	3		
Primary Attributes Used:	13		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	2		
S Vertex Shader	Line: 68	Col: 25	INS

Skip redundant normalization



- Use ***#pragma glsl_no_auto_normalization*** if your normals and tangents don't need normalization

Per-Line Cycle Estimate Total: 79		Emulated Cycle Total: 69	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	69		
Temporary Registers Used:	3		
Primary Attributes Used:	13		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	1		
5 Vertex Shader	Line: 61	Col: 24	INS

Simplify math

- Original expression:

```
fresnel = _ReflectionFactor*((1.0f - _FresnelFactor) + _FresnelFactor*pow(1.0f - RdotN, 5.0f));
```

- Simplified expression:

```
fresnel = _FresnelAdd + _FresnelMul*pow(1.0f - RdotN, 5.0f));
```



Simplify math



- Swap ***pow(x, 5.0f)*** with ***x*x*** (visually effect is similar):

fresnelTemp = 1.0f - RdotN;

fresnel = _FresnelAdd + _FresnelMul*fresnelTemp*fresnelTemp;

Per-Line Cycle Estimate Total: 76		Emulated Cycle Total: 66	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	66		
Temporary Registers Used:	3		
Primary Attributes Used:	13		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	1		
<hr/>			
Vertex Shader	Line: 78	Col: 2	INS

saturate(x) vs. max(x, 0)



- Usually GPU apply ***saturate*** for free...
 - ... but this is not correct for **PowerVR**
 - ***saturate(x) = max(min(x, 1), 0)***
 - so it's more beneficial to use ***max(x, 0)*** for **PowerVR**

Per-Line Cycle Estimate Total: 71		Emulated Cycle Total: 63	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	63		
Temporary Registers Used:	4		
Primary Attributes Used:	13		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	1		
S Vertex Shader	Line: 64	Col: 24	INS

Optimize Point-Matrix multiplication

- *mul(float4x4, float4)* doesn't know anything about actual values
- **PowerVR SGX 540** consume 1 cycle for each scalar *madd* in vertex shader
- So it's beneficial to know that **point = (x, y, z, 1)**
- And some matrices has **(0, 0, 0, 1)** row



Optimize Point-Matrix multiplication

```
inline float3 GetPositionInWorldSpace(float4 positionInLocalSpace)
{
    float3 positionInWorldSpace;

    positionInWorldSpace.x = _Object2World[0].w + dot(_Object2World[0].xyz, positionInLocalSpace.xyz);
    positionInWorldSpace.y = _Object2World[1].w + dot(_Object2World[1].xyz, positionInLocalSpace.xyz);
    positionInWorldSpace.z = _Object2World[2].w + dot(_Object2World[2].xyz, positionInLocalSpace.xyz);

    return positionInWorldSpace;
}
```

```
inline float4 GetPositionInClipSpace(float4 positionInLocalSpace)
{
    float4 positionInClipSpace;

    positionInClipSpace.x = UNITY_MATRIX_MVP[0].w + dot(UNITY_MATRIX_MVP[0].xyz, positionInLocalSpace.xyz);
    positionInClipSpace.y = UNITY_MATRIX_MVP[1].w + dot(UNITY_MATRIX_MVP[1].xyz, positionInLocalSpace.xyz);
    positionInClipSpace.z = UNITY_MATRIX_MVP[2].w + dot(UNITY_MATRIX_MVP[2].xyz, positionInLocalSpace.xyz);
    positionInClipSpace.w = UNITY_MATRIX_MVP[3].w + dot(UNITY_MATRIX_MVP[3].xyz, positionInLocalSpace.xyz);

    return positionInClipSpace;
}
```

Optimize Point-Matrix multiplication

- We see that actual cycle count has grown slightly...
- ... but temporary registers usage reduced

Per-Line Cycle Estimate Total: 70		Emulated Cycle Total: 64	
Compiler:	SGX540 rev 130		
Version:	1.9@2650911		
Emulated Cycles:	64		
Temporary Registers Used:	3		
Primary Attributes Used:	13		
Non-Dependent Texture Loads:	0		
Global USC Instructions:	1		
S Vertex Shader	Line: 97	Col: 31	INS

Fragment shaders optimization



Motivation

- Modern android devices has lots of megapixels 😊
- So every cycle can make you GPU bound



Case study: diffuse lighting



- Per-pixel diffuse lighting from 3 directional light sources
- Target device **HTC One S** (Adreno 220)
- Use **Adreno Profiler** to measure actual cycles spent by GPU



Naive implementation



```
half3 normal = normalize(i.normal);

fixed3 diffuse = global_Ambient.rgb;
diffuse += max(dot(global_LightDirection1.xyz, normal), 0.0f)*global_LightColor1;
diffuse += max(dot(global_LightDirection2.xyz, normal), 0.0f)*global_LightColor2;
diffuse += max(dot(global_LightDirection3.xyz, normal), 0.0f)*global_LightColor3;

fixed4 result;
result.rgb = _Color.rgb*diffuse;
result.a = 1.0f;
return result;
```

Metrics 12.3 Find Redundant Calls Flip Save Save Vertex Data				
#	Render Calls	No Effect	Heavy	Clocks
1	glClear(mask =COLOR DEPTH STENCIL)	0	0	
2	glClear(mask =COLOR DEPTH STENCIL)	4	0	
3	glDrawElements(mode =GL_TRIANGLES, count =2280, t 0		0	339,390.00
4	glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	26,293.00
5	glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	24,826.00





saturate(x) vs. max(x, 0)



```
half3 normal = normalize(i.normal);

fixed3 diffuse = global_Ambient.rgb;
diffuse += saturate(dot(global_LightDirection1.xyz, normal))*global_LightColor1;
diffuse += saturate(dot(global_LightDirection2.xyz, normal))*global_LightColor2;
diffuse += saturate(dot(global_LightDirection3.xyz, normal))*global_LightColor3;

fixed4 result;
result.rgb = _Color.rgb*diffuse;
result.a = 1.0f;
return result;
```


Metrics ▾ 12.3  Find Redundant Calls ▾   Flip  Save Save Vertex Data				
#	Render Calls	No Effect	Heavy	Clocks
1	glClear(mask =COLOR DEPTH STENCIL)	0	0	
2	glClear(mask =COLOR DEPTH STENCIL)	4	0	
3	glDrawElements(mode =GL_TRIANGLES, count =2280, t 0	0	0	266,285.00
4	glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	27,320.00
5	glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	24,842.00

Bake lighting to diffuse cubemap

- Contain lighting environment information about all directional light sources and ambient
- **Normal** vector used to sample diffuse lighting with single *texCUBE* instruction
- Generated on **CPU** with relatively simple code
- Pre-multiplied by 0.5 to increase intensity range
- Size of 32 texels is usually enough



Bake lighting to diffuse cubemap




```
fixed3 diffuse = texCUBE(_DiffuseCubemap, i.normal);  
  
fixed4 result;  
result.rgb = _Color.rgb*diffuse*2.0f;  
result.a = 1.0f;  
return result;
```

Metrics 12.3 Find Redundant Calls Flip Save Save Vertex Data				
#	Render Calls	No Effect	Heavy	Clocks
1	glClear(mask =COLOR DEPTH STENCIL)	0	0	
2	glClear(mask =COLOR DEPTH STENCIL)	4	0	
3	glDrawElements(mode =GL_TRIANGLES, count =2280, t 0	0		155,505.00
4	glDrawElements(mode =GL_TRIANGLES, count =54, typ 9	0		25,045.00
5	glDrawElements(mode =GL_TRIANGLES, count =36, typ 1	0		23,261.00

Case study: specular lighting



- Per-pixel phong specular lighting from 3 directional light sources
- Target device **HTC One S** (Adreno 220)



Naive implementation




```
half3 N = normalize(i.N);
half3 V = normalize(i.V);

half3 R1 = reflect(global_LightDirection1.xyz, N);
half3 R2 = reflect(global_LightDirection2.xyz, N);
half3 R3 = reflect(global_LightDirection3.xyz, N);

fixed3 specular;
specular = pow(saturate( dot(R1, V) ), _shininess)*global_LightColor1;
specular += pow(saturate( dot(R2, V) ), _shininess)*global_LightColor2;
specular += pow(saturate( dot(R3, V) ), _shininess)*global_LightColor3;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =9102, t 0	0	0	559,186.00
glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	26,406.00
glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	24,795.00

Optimize computation scheme



$$L_s = \max(\vec{R} \cdot \vec{L}, 0)^{\text{shininess}} \times \text{LightColor}$$

Optimize computation scheme



```
half3 N = normalize(i.N);
half3 V = normalize(i.V);
half3 R = reflect(V, N);


fixed3 specular;
specular = pow(saturate( dot(R, global_LightDirection1.xyz) ), _shininess)*global_LightColor1;
specular += pow(saturate( dot(R, global_LightDirection2.xyz) ), _shininess)*global_LightColor2;
specular += pow(saturate( dot(R, global_LightDirection3.xyz) ), _shininess)*global_LightColor3;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =9102, t 0		0	479,037.00
glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	27,246.00
glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	24,740.00

Use fast pow approximation



- **pow** is usually scalar and consume a lot of cycles
- approximation can make up to 4 pows in just 2 cycles!!



Use fast pow approximation



```
half3 N = normalize(i.N);
half3 V = normalize(i.V);
half3 R = reflect(V, N);

half3 specularVector;
specularVector.x = saturate( dot(R, global_LightDirection1.xyz) );
specularVector.y = saturate( dot(R, global_LightDirection2.xyz) );
specularVector.z = saturate( dot(R, global_LightDirection3.xyz) );

specularVector = saturate(specularVector*_specularPower + (1.0f - _specularPower));
specularVector = specularVector*specularVector;

fixed3 specular;
specular = specularVector.x*global_LightColor1;
specular += specularVector.y*global_LightColor2;
specular += specularVector.z*global_LightColor3;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =9102, t 0	0	0	425,891.00
glDrawElements(mode =GL_TRIANGLES, count =54, typ 9	0	0	26,111.00
glDrawElements(mode =GL_TRIANGLES, count =36, typ 1	0	0	24,606.00

Move calculations to vertex shader

```
half3 R = normalize(i.R);

half3 specularVector;
specularVector.x = saturate( dot(R, global_LightDirection1.xyz) );
specularVector.y = saturate( dot(R, global_LightDirection2.xyz) );
specularVector.z = saturate( dot(R, global_LightDirection3.xyz) );

specularVector = saturate(specularVector*_specularPower + (1.0f - _specularPower));
specularVector = specularVector*specularVector;

fixed3 specular;
specular = specularVector.x*global_LightColor1;
specular += specularVector.y*global_LightColor2;
specular += specularVector.z*global_LightColor3;
```


Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =9102, t 0		0	320,863.00
glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	26,767.00
glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	24,737.00

Bake lighting to specular cubemap

- **Reflected view** vector used to sample specular lighting with single *texCUBE* instruction
 - Size of 64 texels is usually enough
 - Pre-multiplied by 0.5 to increase intensity range
-
- Really interesting BRDF can be baked using this approach



Bake lighting to specular cubemap





```
return texCUBE(_SpecularCubemap, i.R)*2.0f;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =9102, t 0		0	160,711.00
glDrawElements(mode =GL_TRIANGLES, count =54, typ 9		0	25,567.00
glDrawElements(mode =GL_TRIANGLES, count =36, typ 1		0	22,984.00

Case study: specular map



- Specular intensity map greatly improve realism
- Optimize for **HTC One S**




Naive approach

- Store in alpha channel of diffuse texture
- Use ARGB32 texture format

Color



Alpha



Naive approach



```
fixed3 specular = texCUBE(_SpecularCubemap, i.V);
fixed4 mainTex = tex2D(_MainTex, i.uv);

fixed4 result;

result.rgb = mainTex.rgb + specular*mainTex.a;
result.a = 1.0f;

return result;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =36, type 1	0	0	686,881.00

Use better texture format



- Adreno GPUs support compressed texture format with 8 bits per pixel which contains alpha channel
- The problem is you should select different format for all platforms
- **Mali Utgard** doesn't support compressed textures with alpha channel 😞

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =36, type 1	1	0	483,893.00

Recover specular map in shader

- Use ETC1 to store diffuse map
- Compute diffuse luminance for each pixel

$$\text{luminance} = 0.3 * R + 0.58 * G + 0.12 * B$$

- Scale and bias this luminance to get specular intensity

$$\text{specIntensity} = \text{saturate}(\text{luminance} * \text{scale} + \text{bias})$$

- Combine both formulas

$$\text{specIntensity} = \text{saturate}(\text{dot}(\text{RGB1}, \text{specConst}))$$

$$\text{specConst} = (0.3 * \text{scale}, 0.58 * \text{scale}, 0.12 * \text{scale}, \text{bias})$$



Recover specular map in shader unity

```
fixed3 specular = texCUBE(_SpecularCubemap, i.V);
fixed4 mainTex = tex2D(_MainTex, i.uv);
fixed specularIntensity = saturate(dot(mainTex, i.specularConst));

fixed4 result;

result.rgb = mainTex.rgb + specular*specularIntensity;
result.a = 1.0f;

return result;
```

Render Calls	No Effect	Heavy	Clocks
glClear(mask =COLOR DEPTH STENCIL)	0	0	
glClear(mask =COLOR DEPTH STENCIL)	4	0	
glDrawElements(mode =GL_TRIANGLES, count =36, type 1	0	0	465,042.00

Questions ?

